

Summary

- Model Parameters → Estimated from the data
- Model Hyperparameters → Can't be estimated from the data

Model hyperparameters are often referred to as parameters because they are the parts of the machine learning that must be set manually and tuned.

5 Epochs, Batches, Batch Sizes & Iterations

- Need these terms only if the dataset is **large**
- Break down the dataset into smaller "**chunks**" and feed those chunks to the neural network one-by-one

Epochs

When the ENTIRE dataset is passed forward and backward through the neural network only ONCE

We use multiple epochs to help our model generalize better

Batch & Batch Size

We divide large datasets into smaller **batches** and feed those batches into the neural network

Batch Size: Total number of training examples in a Batch

Neural Network Architectures

Iterations

Number of batches needed to complete one epoch

Alternatively,

Number of
batches

=

Number of
iterations for
one epoch

Search



Suppose you have a dataset of **34000** training examples
and you divide the dataset into batches of **500**.

To complete 1 epoch, it would have taken **68** iterations

Conclusion to Terms Used in NNs

How many hidden layers should I choose?

Which Activation function to pick?

**Experiment, Experiment,
Experiment!**

Types of Learning

3 Main Types:

1. **Supervised** Learning
2. **Unsupervised** Learning
3. **Reinforcement** Learning

1 Supervised Learning

- Algorithms designed to learn by example
- Models are trained on **well-labelled** data

Each example is a pair consisting of:

- **Input** Object (typically a vector)
- **Desired** Output (Supervisory Signal)

During training,

SL algorithm searches for patterns that correlate with the desired output

After training,

takes in unseen inputs and determine which label to classify it to

Objective of a Supervised Learning Model:

**Predict the correct label for
unseen data**

$$y = f(x)$$

x = Input

y = Predicted Output

Supervised Learning



Classification

Regression

1 Classification

Take input data and assign it to a class/category

Example: Is email **spam** or **not spam**

Models finds features in the data that correlate to a class and creates a **mapping function**

This mapping function will be used to classify unseen data

Search



Popular Classification Algorithms

1. Linear Classifiers
2. Support Vector Machines
3. K-Nearest Neighbour
4. Random Forest

2 Regression

Model attempts to find a relationship between dependent & independent variables

Goal: To predict continuous values such as Test Scores

Equation for Basic Linear Regression

$$\hat{y} = w[0]*x[0] + w[1]*x[1] + \dots + w[i]*x[i] + b$$

$x[i]$ = i^{th} Input Feature

$w[i]$, b = Parameters

Many Input Features

Popular Regression Algorithms

1. Linear Regression
2. Lasso Regression
3. Multivariate Regression

Applications of Supervised Learning

1. Bioinformatics
2. Object Recognition
3. Spam Detection
4. Speech Recognition

2 Unsupervised Learning

- Uses to manifest **underlying patterns** in data
- Used in exploratory data analysis
- Does not use labelled data, rather relies on the **data features**
- **Goal:** Analyze data and find important underlying patterns


2 Unsupervised Learning

- Uses to manifest **underlying patterns** in data
- Used in exploratory data analysis
- Does not use labelled data, rather relies on the **data features**
- **Goal:** Analyze data and find important underlying patterns


1 Clustering

Process of grouping data into different **clusters or groups**

Goal: To predict continuous values such as Test Scores



Good Clustering



Bad Clustering

Clustering:

1. Partitional Clustering

Each data point can belong to a single cluster

2. Hierarchical Clustering

Clusters within clusters

Data point may belong to many clusters

Popular Clustering Algorithms

1. K-Means
2. Expectation Maximization
3. Hierarchical Cluster Analysis (HCA)

2 Association

Attempts to find relationships between different entities

Example: Market Basket Analysis

Applications of **Unsupervised Learning**

1. AirBnb
2. Amazon
3. Credit Card Fraud Detection

③ Reinforcement Learning

- Enables an **agent** to learn in an interactive **environment** by trial & error based on feedback from its own actions & experiences
- Uses **rewards** & **punishments** as signals for positive & negative behaviour

3 Reinforcement Learning


- Enables an **agent** to learn in an interactive **environment** by trial & error based on feedback from its own actions & experiences
- Uses **rewards** & **punishments** as signals for positive & negative behaviour
- **Goal:** Find a suitable model that would maximize the total cumulative reward

- Maximize the points won in a game over many moves

Penalized when they make wrong decisions

Rewarded when they make the right ones

Usually modelled as a **Markov Decision Process**



Applications of Reinforcement Learning

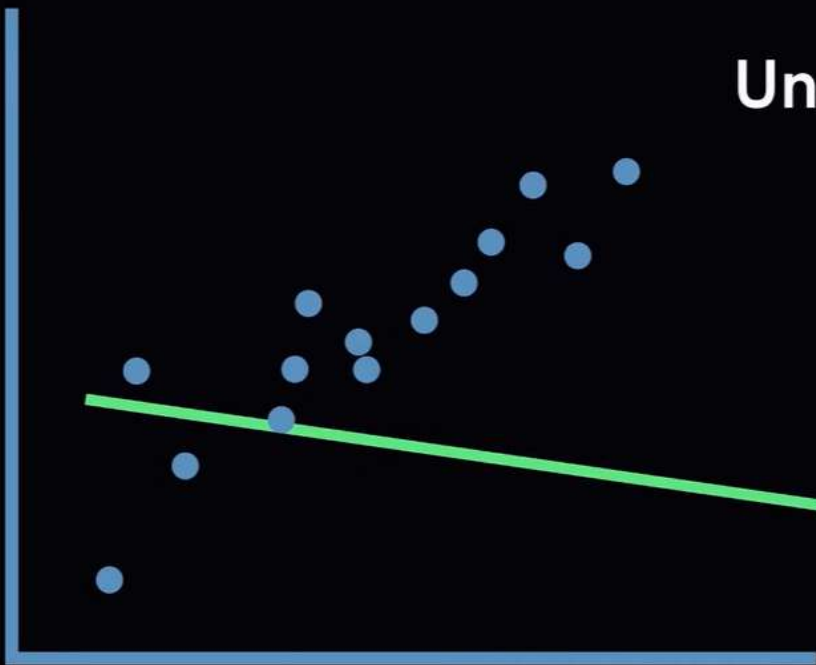
1. Robotics
2. Business Strategy Planning
3. Traffic Light Control
4. Web System Configuration

Core Problem in Deep Learning

Model should perform well on training data **AND** new test data

Most common problem faced is **Overfitting**

Underfitting








Not the best, but will perform better on training data **and** testing data

Tackling Overfitting

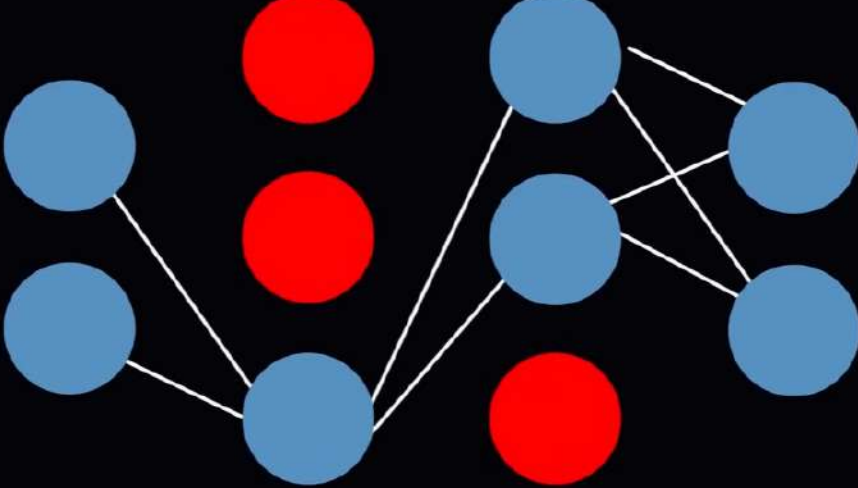
Tackling Overfitting

1. Dropout
2. Augmentation
3. Early Stopping

1 Dropout



Dropout randomly removes some nodes & their connections



The diagram shows a network of 10 blue circular nodes and 10 white lines representing connections. Three red circular nodes are shown, representing nodes that have been randomly removed (dropped out) from the network. The connections between the remaining blue nodes are preserved, illustrating how the network structure changes during dropout.

Search



② Dataset Augmentation

More Data  Better Model

Fake Data

2 Dataset Augmentation

Apply transformations on the existing dataset to get synthesize more data

3 Early Stopping

Training error **decreases** steadily

But, validation error **increases** after a certain point

Neural Network Architectures

Fully-Connected Feed Forward Neural Networks

Each neuron is connected to every subsequent layer with no backward connections

Neurons have activation functions

- Linear
- Sigmoid
- Tanh
- ReLU



Non-Linear

Inputs

Outputs

Hidden Layers

Neurons per hidden layer

Activation functions

**More
Neurons**



**Larger
Computational
Resources**

Recurrent Neural Networks

Feed-Forward Neural Networks

take in fixed-sized inputs

returns fixed-sized outputs

Feed-Forward Neural Networks

take in fixed-sized inputs
returns fixed-sized outputs

Can't model every problem today

**How does this apply to
Neural Networks?**

Information about the **past must be supplied**

Vanilla Neural Networks can't handle **sequential data**

Sequential Data is data in a sequence

Sequential Data is data in a sequence

**Vanilla NNs don't share parameters
across time**

It was raining on
Tuesday

On Tuesday, it was
raining



Mean the same thing

The diagram consists of two green arrows pointing upwards from a central point below the text 'Mean the same thing' towards the two sentences above. The left arrow points to 'It was raining on Tuesday' and the right arrow points to 'On Tuesday, it was raining'.

Sharing parameters gives the network the ability to look for a given feature **everywhere in the sequence**, rather than in just a certain area

Deal with **variable length** sequences


Maintain sequence **order**


Keep track of **long-term dependencies**

Share parameters across the sequence

Recurrent Neural Networks

Uses a **feedback loop** in the hidden layers





Training an RNN

Uses the Backpropagation algorithm

Backprop applied for every **sequence data point**

Backpropagation through Time (BTT)

Training an RNN

Uses the Backpropagation algorithm


Backprop applied for every **sequence data point**

Backpropagation through Time (BTT)

Intelligent Predicting

Predict letters based on previously-typed letters

All letters typed are **equally important** in prediction!



Vanishing Gradient Problem

Short-term memory of an RNN is due to VGP

Due to the nature of Backpropagation

In Backpropagation:

$$\text{Error} = \left(\begin{array}{c} \text{Predicted} \\ \text{Output} \end{array} - \begin{array}{c} \text{Expected} \\ \text{Output} \end{array} \right)^2$$

In Backpropagation:

$$\text{Error} = \left(\begin{array}{c} \text{Predicted} \\ \text{Output} \end{array} - \begin{array}{c} \text{Expected} \\ \text{Output} \end{array} \right)^2$$

$$W = w + n \frac{de}{dw}$$

Gradients of a layer are calculated based on the gradients of the **previous** layer

Gradients of a layer are calculated based on the gradients of the **previous** layer

If initial gradient is small, adjustments to the subsequent layers will be **smaller** giving rise to vanishing gradients

Every time step in RNN = A Layer

To train, we use **BTT**

Gradients used make adjustments to weights and biases

Because of **VGP**, RNNs are unable to learn long-range dependencies

"It was raining on Tuesday"

"It" and "was" may not be considered

Model will have to guess based on "on Tuesday"

LSTM - Long Short Term Memory

GRNN - Gated RNN

Capable of learning long-term dependencies
using **gates**

GRNNs

- Update Gate
- Reset Gate

LSTMs

- Update Gate
- Reset Gate
- Forget Gate

Applications of RNNs

1. Natural Language Processing
2. Sentiment Analysis
3. DNA Sequence Classification
4. Speech Recognition
5. Language Translation

Convolutional NNs

- Inspired by the organization of neurons in the **visual cortex** of the human brain
- Good for processing data like **images**, **audio** and **video**